Submission Checklist — CISC7021 Group 26 CALM Project

Owner: FONG HIO LAM (MC654953) **Package:** weekly deliverables W0–W6 **Last verified:** 8 November 2026

1. Pre-Submission Gate

Tick every box before packaging. Each row names the check and how it was run at the stated date, so it can be reproduced.

#	Check	Command / Method	Result
1	Every week has a submission directory	`ls project/output/w*_fonghiolam`	Pass — `w0`–`w6` all present
2	Working copy and submission copy agree	byte comparison of `project/wX` against `project/output/wX_fonghiolam`	Pass
3	All Python compiles	`py_compile` over all `.py` under `output/`	Pass — 27/27
4	Smoke tests pass	W1 `--mock --limit 3`; W3 `side_memory.py`, `conflict_gate.py`; W4 `--mock`; W5 `--mock`	Pass
5	English only	recursive CJK scan of `.md/.py/.json/.txt/.csv/.yaml/.html/.tex`	Pass — 0 CJK in authored text; one retained model-output exception, see §5
6	No vendor labelling metadata	scan for the machine-generation labelling keys used by document exporters	Pass — 0 hits
7	No disclaimer lines	scan for sentences claiming machine generation or reference-only status	Pass — 0 hits
8	No byte-compiled caches shipped	scan for cache directories	Pass — removed, W1 archive rebuilt
9	Every README quotes the metric convention it uses	manual read of W1–W5 READMEs	Pass
10	Git history clean and pushed	`git status`, `git log`, `git ls-remote` in `project_repo`	Pass — every W0–W6 commit is on `origin/main`, working tree clean
11	W3 memory channel meets the six Week-6 handbook checks	handbook S1–S6 against `side_memory.py`, `real_w3_components.py`, `w3_memory_channel_gpt2xl.json`	Pass — key is `(subject, relation)`, answer in the value, read path before generation, recall split 0.80 / 0.45 on n = 20, shuffled control 0.00, no 1.00 quoted as accuracy

2. Per-Week File Manifest

Week 0 — Literature Review & Taxonomy

- [x] `output/w0_fonghiolam/Section2_3_Literature_Review_Taxonomy_FongHioLam_W0.md` (178 lines)
- [x] Mirrored at `project/w0/FONG HIO LAM/Section2_3_Literature_Review_Taxonomy_FongHioLam_W0.md`
- [x] Reference list present, `\citep{}` keys mapped to the bibliography table

Week 1 — Data Loaders & Evaluation Metrics

- [x] `evaluate_pipeline.py` (259 lines) — mock, pre-scored and real-model modes
- [x] `metrics.py` (230 lines) — comprehensive metric battery
- [x] `real_eval_w1.py` (149 lines)
- [x] `w1_alignment_decision.md` (105 lines) — model and dataset decisions
- [x] `loaders/` (`__init__`, `base`, `counterfact`, `zsre`, `mquake`)
- [x] `sample_data/` — CounterFact, ZsRE, MQuAKE samples
- [x] Real runs: `w1_counterfact_gpt2xl.json` (100 cases), `w1_zsre_gpt2xl.json` (100 cases)
- [x] `README.md` (159 lines)

Week 2 — MEMIT & AlphaEdit

- [x] `memit_repro.py` (543 lines), `alphaedit_repro.py` (307 lines)
- [x] `real_eval_w2.py` (489 lines), `real_w2_edit.py` (192 lines), `rescore_w2.py` (76 lines)
- [x] `smoke_outputs/memit_smoke.json`, `smoke_outputs/alphaedit_smoke.json`
- [x] Real runs: `w2_xl_memit.json`, `w2_xl_alphaedit.json` (12 cases, 3 layers)
- [x] `README.md` (217 lines) — states exact vs containment vs early-window rewrite

Week 3 — SideMemory & ConflictGate

- [x] `side_memory.py` (266 lines), `conflict_gate.py` (399 lines)
- [x] `real_w3_components.py` (455 lines) — 100-request run driver
- [x] `w3_memory_channel.py` (218 lines) — answer-path recall study, cue-only baseline included
- [x] `probe_sweep.py` (174 lines), `probe_likelihood.py` (221 lines), `probe_sampling.py` (168 lines)
- [x] `probe_false_reject.py` (175 lines) — false-reject cost of the sampling variant
- [x] `canary_strata.py` (131 lines) — stratum join
- [x] Real runs: `w3_components_gpt2xl.json`, `probe_sweep_gpt2xl_v2.json`, `probe_likelihood_gpt2xl.json`, `probe_sampling_gpt2xl.json`, `probe_false_reject_gpt2xl.json`, `canary_strata_gpt2xl.json`, `w3_memory_channel_gpt2xl.json`
- [x] `README.md` (503 lines) — three canary denominators, the sampling cost table, and the memory-channel controls
- [x] Mirrored to the team repository (`side_memory_router`)

Week 4 — MQuAKE & WikiBigEdit

- [x] `mquake_eval.py` (435 lines), `wikibigedit_eval.py` (393 lines)
- [x] `real_w4_mquake.py` (291 lines), `real_w4_wikibigedit.py` (290 lines), `diag_w4_write.py` (70 lines)
- [x] Real runs: `mquake_results.json` (50 samples), `wikibigedit_results.json` (50 samples)
- [x] `README.md` (427 lines)

Week 5 — Conflict-Type Error Analysis

- [x] `conflict_analysis.py` (547 lines)
- [x] `conflict_analysis_report.txt` (WikiBigEdit), `conflict_analysis_report_mquake.txt`
- [x] Six figures, all generated from the two result files, all text English
- [x] `README.md` (265 lines)

Week 6 — Cross-Review & Submission

- [x] `Cross-Review_Report.md` — inventory, per-week review, compliance, open items
- [x] `Submission_Checklist.md` — this file

---

3. Numbers to Quote in the Presentation

Metric	Value	Source file
W1 CounterFact rewrite (no editing)	0.01	`w1_counterfact_gpt2xl.json`
W2 MEMIT rewrite (containment / exact)	1.00 / 0.00	`w2_xl_memit.json`
W2 AlphaEdit paraphrase	0.708	`w2_xl_alphaedit.json`
W3 decision split (WRITE / ROUTE / REJECT)	97 / 1 / 2	`w3_components_gpt2xl.json`
W3 memory recall, canonical cue	16/20 = 0.80	`w3_memory_channel_gpt2xl.json`
W3 memory recall, paraphrase cue	9/20 = 0.45	`w3_memory_channel_gpt2xl.json`
W3 memory recall, no note (baseline)	0/20 = 0.00	`w3_memory_channel_gpt2xl.json`
W3 memory recall, shuffled key (control)	0/20 = 0.00	`w3_memory_channel_gpt2xl.json`
W3 canary recall, raw	11/20 = 0.55	`w3_components_gpt2xl.json`
W3 canary recall, model-known stratum	11/14 = 0.786	`canary_strata_gpt2xl.json`
W3 canary recall, sampling upper bound	14/14 = 1.00	`canary_strata_gpt2xl.json`
W3 false rejects, deployed gate	2/100 = 0.02	`probe_false_reject_gpt2xl.json`
W3 false rejects, sampling-augmented gate	14/100 = 0.14	`probe_false_reject_gpt2xl.json`
W4 MQuAKE rewrite / generation	0.90 / 0.18	`mquake_results.json`
W4 WikiBigEdit rewrite / locality	0.82 / 0.24	`wikibigedit_results.json`
W5 WikiBigEdit error rate	36.0%	`conflict_analysis_report.txt`
W5 MQuAKE error rate	94.0%	`conflict_analysis_report_mquake.txt`

---

4. Re-Run Commands

All commands assume the project root as working directory and the project virtual environment on `PATH`.

```
# W1
python output/w1_fonghiolam/evaluate_pipeline.py --mock --limit 3

# W2
python output/w2_fonghiolam/memit_repro.py
python output/w2_fonghiolam/alphaedit_repro.py

# W3 (module self-tests)
python output/w3_fonghiolam/side_memory.py
python output/w3_fonghiolam/conflict_gate.py
# W3 (studies, need the gpt2-xl checkpoint)
python output/w3_fonghiolam/real_w3_components.py --out output/w3_fonghiolam/w3_components_gpt2xl.json
python output/w3_fonghiolam/canary_strata.py --run w3_components_gpt2xl.json \
    --sampling probe_sampling_gpt2xl.json --out canary_strata_gpt2xl.json

# W4
python output/w4_fonghiolam/mquake_eval.py --mock
python output/w4_fonghiolam/wikibigedit_eval.py --mock

# W5
python output/w5_fonghiolam/conflict_analysis.py --mock
```

---

5. Packaging

- ☑ Seven week directories, `w0_fonghiolam` through `w6_fonghiolam`
- ☑ One `README.md` per week, each with its own result tables inlined
- ☑ No byte-compiled cache directories, no compiled `.pyc` files, no editor backup files
- ☑ No Chinese text, no AI metadata, no disclaimers
- ☑ PDF export of the W6 deliverables — both Markdown files converted with `pandoc` (GFM to HTML5) and printed by headless Edge to A4: `CISC7021_Group26_W6_Cross-Review_Report.pdf` (10 pages) and `CISC7021_Group26_W6_Submission_Checklist.pdf` (4 pages). Browser header and footer suppressed; the exported files were re-scanned afterwards and show 0 CJK characters in the extracted text, 0 labelling keys in the file bytes, and no XMP packet. Both PDFs are tracked in the team repository through an exception in `.gitignore`, which otherwise ignores `*.pdf` as a build artefact
- ☑ Final push to the team repository — all W0–W6 files plus the two W6 PDFs tracked, `origin/main` in sync, working tree clean
- ☑ Known exception recorded: `probe_false_reject_gpt2xl.json` keeps three CJK fragments (lines 340, 5299, 6443) inside model-generated text, retained as run evidence and excluded from the authored-text scan